

---

# CS 4375 – Introduction to Machine Learning

---

Evaluating Machine Learning Models

Erick Parolin



# Model Evaluation

- **So far, we covered a plenty of learners for classification task...**
  - Learning algorithm
  - How to use them for classification
  - How they explore the feature space
  - Etc.
- **Now we will learn how to evaluate the *performance* of these models**
  - **Focus on their predictive capability (NOT time/space complexity)**
  - How to evaluate the performance of a model.
  - How to obtain reliable estimates.
  - How to compare performance of the different models.

# Metrics for Performance Evaluation

## Confusion Matrix

Confusion Matrix		Predicted Class	
Actual Class	Positive (Class=1)	Positive (Class=1)	Negative (Class=0)
	Negative (Class=0)	FP	TN
TP		FN	

Each record from the dataset used for model evaluation will fall into one of the cells of the confusion matrix:

TP: True Positive

FP: False Positive

FN: False Negative

TN: True Negative

# Metrics for Performance Evaluation

## Accuracy

Confusion Matrix		Predicted Class	
		Positive (Class=1)	Negative (Class=0)
Actual Class	Positive (Class=1)	TP	FN
	Negative (Class=0)	FP	TN

**TP:** True Positive  
**FP:** False Positive  
**FN:** False Negative  
**TN:** True Negative

$$\text{Accuracy} = \frac{TP + TN}{TP + FP + FN + TN}$$

# Metrics for Performance Evaluation

## Accuracy

- **Limitation:** what if I have *unbalanced dataset*?
- **Example:**
  - Number of Class 1 instances is 99,900
  - Number of Class 0 instances is 100
  - If the model is a constant that predicts everything as Class 1, we have accuracy 99.9%
  - Accuracy may be misleading: the model does not detect any Class 0!!!

# Metrics for Performance Evaluation

## Cost Matrix

Cost Matrix		Predicted Class	
		Positive (Class=1)	Negative (Class=0)
Actual Class	Positive (Class=1)	C(1 1)	C(0 1)
	Negative (Class=0)	C(1 0)	C(0 0)

$C(P | A)$ : cost of predicting class P when actual class is A

# Metrics for Performance Evaluation

## Cost of Classification – Numerical Example

### Dataset

**Class 1:** 357 examples

**Class 0:** 5,303 examples

Model 1		Predicted Class	
		Positive (Class=1)	Negative (Class=0)
Actual Class	Positive (Class=1)	282	75
	Negative (Class=0)	462	4,841

**Model 1**  
**Accuracy:** 90.5%  
**Cost:** 23,958

### Cost Matrix

Cost Matrix		Predicted Class	
		Positive (Class=1)	Negative (Class=0)
Actual Class	Positive (Class=1)	-1	200
	Negative (Class=0)	20	0

Model 2		Predicted Class	
		Positive (Class=1)	Negative (Class=0)
Actual Class	Positive (Class=1)	332	25
	Negative (Class=0)	607	4,696

**Model 2**  
**Accuracy:** 88.8%  
**Cost:** 16,808

# Metrics for Performance Evaluation

**Precision:** measures how well a model can correctly predict positive instances

Confusion Matrix		Predicted Class	
		Positive (Class=1)	Negative (Class=0)
Actual Class	Positive (Class=1)	TP	FN
	Negative (Class=0)	FP	TN

TP: True Positive

FP: False Positive

FN: False Negative

TN: True Negative

$$\textit{Precision} = \frac{TP}{TP + FP}$$

# Metrics for Performance Evaluation

**Recall:** measures how well a model can identify positive instances

Confusion Matrix		Predicted Class	
		Positive (Class=1)	Negative (Class=0)
Actual Class	Positive (Class=1)	TP	FN
	Negative (Class=0)	FP	TN

**TP:** True Positive  
**FP:** False Positive  
**FN:** False Negative  
**TN:** True Negative

$$\text{Recall} = \frac{TP}{TP + FN}$$

# Metrics for Performance Evaluation

**F-score:** a measure of the harmonic mean of precision and recall.

Confusion Matrix		Predicted Class	
		Positive (Class=1)	Negative (Class=0)
Actual Class	Positive (Class=1)	TP	FN
	Negative (Class=0)	FP	TN

TP: True Positive  
FP: False Positive  
FN: False Negative  
TN: True Negative

$$F_\beta = \frac{(\beta^2 + 1) \cdot \text{Prec} \cdot \text{Rec}}{\beta^2 \cdot \text{Prec} + \text{Rec}}$$

Most common ( $\beta = 1$ ):

$$F1 = 2 \frac{\text{Prec} \cdot \text{Rec}}{\text{Prec} + \text{Rec}} = \frac{2\text{TP}}{2\text{TP} + \text{FP} + \text{FN}}$$

# What if we have more than two classes?

## Multiclass (classes are mutually exclusive)

- $Prec_k$ : Measures how many of the predicted positive instances for class  $k$  are correct.
- $Rec_k$ : Measures how many actual positive instances for class  $k$  were correctly predicted.
- $F1_k$ : The harmonic mean of precision and recall for class  $k$ .
- **Accuracy** is computed in the same manner.

$$Prec_k = \frac{TP_k}{TP_k + FP_k}$$

$$Rec_k = \frac{TP_k}{TP_k + FN_k}$$

$$F1_k = 2 \frac{Prec_k \cdot Rec_k}{Prec_k + Rec_k}$$

**Model level performance** can be computed by averaging the measurements across all classes  $k \in K$ .

- **Micro average**: Calculates metrics globally by counting the total TP, FP and FN.
- **Macro average**: Calculates metrics for each label and find the unweighted mean (regardless the proportion of each class).
- **Weighted average**: Calculate metrics for each label and find the average weighted by support (# of true examples for each class).

# What if we have more than two classes?

## Multiclass Example

Confusion Matrix		Predicted Class			Support
		Class=0	Class=1	Class=2	
Actual Class	Class=0	512	12	22	546
	Class=1	2	77	13	92
	Class=2	36	59	831	926

$$\text{Acc} = \frac{512 + 77 + 831}{546 + 92 + 926} = 90.8\%$$

$$\text{Rec}_0 = \frac{512}{546} = 93.8\%$$

$$\text{Rec}_1 = \frac{77}{92} = 83.7\%$$

$$\text{Rec}_2 = \frac{831}{926} = 89.7\%$$

$$\text{Prec}_0 = \frac{512}{550} = 93.1\%$$

$$\text{Prec}_1 = \frac{77}{148} = 52.0\%$$

$$\text{Prec}_2 = \frac{831}{866} = 96.0\%$$



$$\text{Prec}_{\text{macro}} = \frac{93.1\% + 52.0\% + 96.0\%}{3} = 80.4\%$$

$$\text{Prec}_{\text{micro}} = \frac{512 + 77 + 831}{550 + 148 + 866} = 90.8\% = \text{Acc}$$

$$\text{Prec}_{\text{weighted}} = \frac{93.1\% * 546 + 52\% * 92 + 96\% * 926}{546 + 92 + 926} = 92.4\%$$

$$\text{Rec}_{\text{macro}} = \frac{1}{K} \sum_{k \in K} \text{Rec}_k = \frac{93.8\% + 83.7\% + 89.7\%}{3} = 89.1\%$$

$$\text{Rec}_{\text{micro}} = \frac{\sum_{k \in K} \text{TP}_k}{\sum_{k \in K} \text{TP}_k + \text{FN}_k} = \frac{512 + 77 + 831}{546 + 92 + 926} = 90.8\% = \text{Acc}$$

$$\text{Rec}_{\text{weighted}} = \frac{\sum_{k \in K} \text{Rec}_k * \text{Sup}_k}{N} = \frac{93.8\% * 546 + 83.7\% * 92 + 89.7 * 926}{546 + 92 + 926} = 90.8\% = \text{Acc}$$

# What if we have more than two classes?

**Multi-label** (each example can be assigned more than one label)

$$\text{Accuracy} = \frac{1}{N} \sum_{i=1}^N \frac{|Y_i \cap \hat{Y}_i|}{|Y_i \cup \hat{Y}_i|}$$

Where  $Y_i \in \{0,1\}^{|K|}$  is the true label set (the set of true labels) and  $\hat{Y}_i \in \{0,1\}^{|K|}$  is the set of predicted labels for instance  $i$ . It measures how much the predicted and true labels overlap.

$$\text{Recall} = \frac{1}{N} \sum_{i=1}^N \frac{|Y_i \cap \hat{Y}_i|}{|Y_i|}$$

Measures the fraction of relevant labels that were predicted correctly.

$$\text{Precision} = \frac{1}{N} \sum_{i=1}^N \frac{|Y_i \cap \hat{Y}_i|}{|\hat{Y}_i|}$$

Measures the fraction of relevant labels among the predicted labels.

# What if we have more than two classes?

**Multi-label** (each example can be assigned more than one label)

$$F1 = \frac{1}{N} \sum_{i=1}^N \frac{2 \times |Y_i \cap \hat{Y}_i|}{|Y_i| + |\hat{Y}_i|}$$

Harmonic mean of precision and recall over all instances.

**Example:**

$$Y = \{\{0,1,1\}, \{1,0,1\}, \{1,1,1\}, \{1,0,1\}\}$$

$$\hat{Y} = \{\{0,0,1\}, \{1,0,1\}, \{1,1,0\}, \{1,1,0\}\}$$

$$F1 = \frac{1}{4} \cdot \left[ \frac{2 \cdot 1}{3} + \frac{2 \cdot 2}{4} + \frac{2 \cdot 2}{5} + \frac{2 \cdot 1}{4} \right] = 74.2\%$$

# Evaluation Metrics – Summary

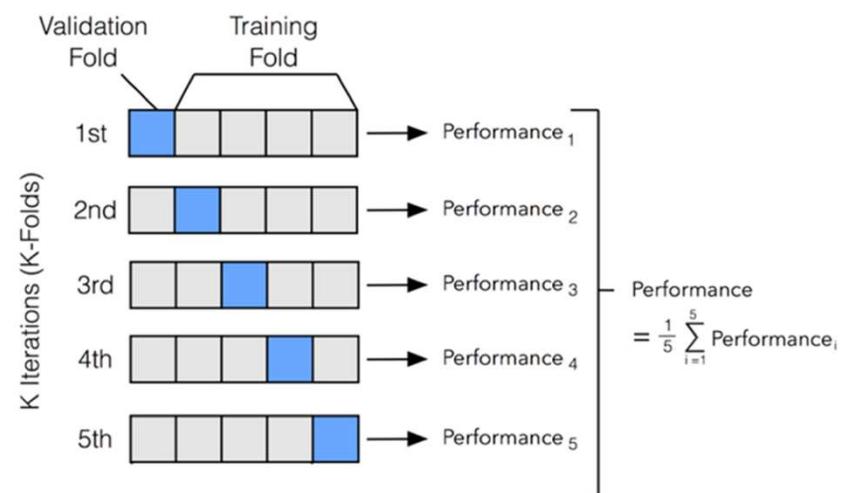
## Other Evaluation Metrics

Confusion Matrix		Predicted Class		<i>Recall or Sensitivity or True Positive Rate</i> $\frac{TP}{TP + FN}$	<i>False Negative Rate</i> $\frac{FN}{TP + FN}$
		Positive (Class=1)	Negative (Class=0)		
Actual Class	Positive (Class=1)	TP	FN		
	Negative (Class=0)	FP	TN	<i>Specificity or True Negative Rate</i> $\frac{TN}{FP + TN}$	<i>False Positive Rate</i> $\frac{FP}{FP + TN}$
$\frac{TP}{TP + FP}$		<i>Precision</i> $\frac{TP}{TP + FP}$	<i>Negative Predictive Value</i> $\frac{TN}{FN + TN}$	<i>Accuracy</i> $\frac{TP + TN}{TP + FP + FN + TN}$	

# Model Validation

## Model Validation and Resampling Methods

- **Cross-validation:** A resampling technique used to evaluate model performance by partitioning the data into multiple subsets (folds) and rotating the training and testing phases.
- Cross-validation divides the data into  $k$  folds.
- The model is trained on  $k-1$  folds and tested on the remaining fold.
- This process is repeated  $k$  times, and the final performance is the average across all folds.



# Model Validation

## Model Validation and Resampling Methods

- **Hold-out:** A simple validation method where the dataset is split into two or more distinct sets (typically training and test sets) to evaluate the model's performance.
- **Bootstrap:** A statistical resampling method used to estimate the accuracy of a model by generating multiple training datasets through sampling with replacement from the original dataset.

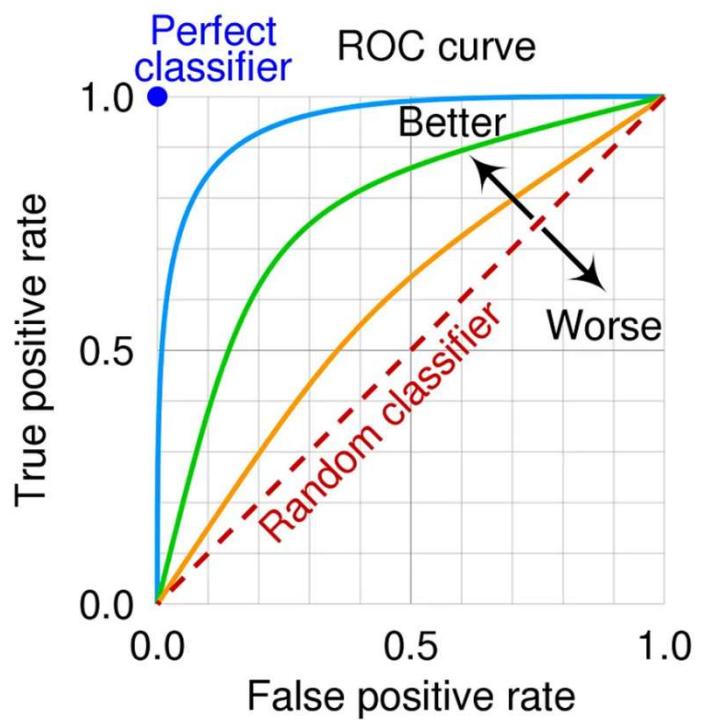
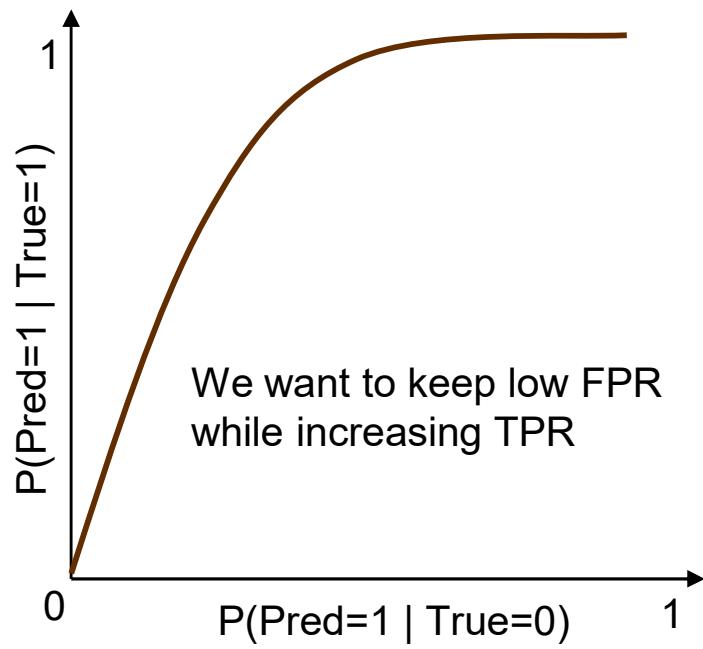
# ROC

## Receiver Operating Characteristics (ROC)

- **Threshold Independence:** ROC shows how the model performs across all possible decision thresholds (accuracy, precision and recall depend on a specific threshold).
- **Trade-offs:** ROC curve highlights the trade-off between True Positive Rate (Recall) and False Positive Rate. This helps understand how the model balances detecting true positives while avoiding false positives.
  - Useful when getting a positive wrong costs more than getting a negative wrong (or vice versa).
- **Class Imbalance:** Useful in cases of class imbalance, where accuracy can be misleading. It provides a clearer picture of how well the model discriminates between classes, even when one class is much more frequent.

# ROC

## Receiver Operating Characteristics (ROC)



# ROC

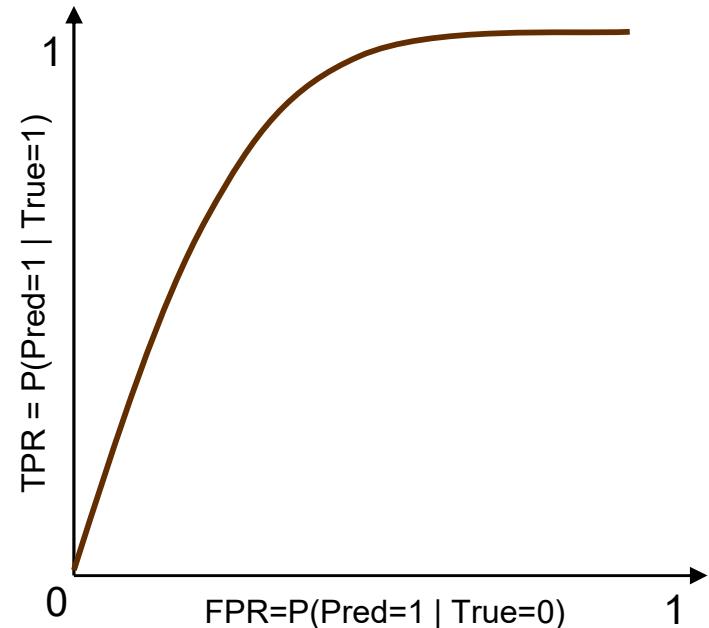
## Algorithm for Creating ROC Curve

**Step 1:** Sort test set by predictions  $P(\text{Class}=1)$  in decreasing order.

**Step 2:** Set thresholds along these examples (it can be based on equal area/width bins or example by example)

**Step 3:** Compute TPR & FPR for each threshold of Step 2

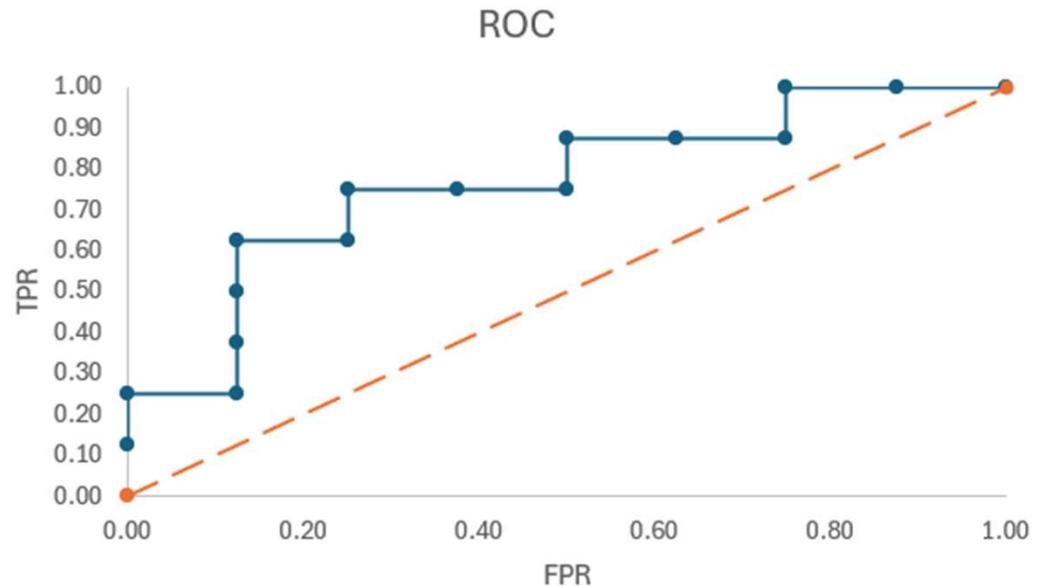
**Step 4:** Connect the dots in the plot



# ROC

## Example of ROC Plot

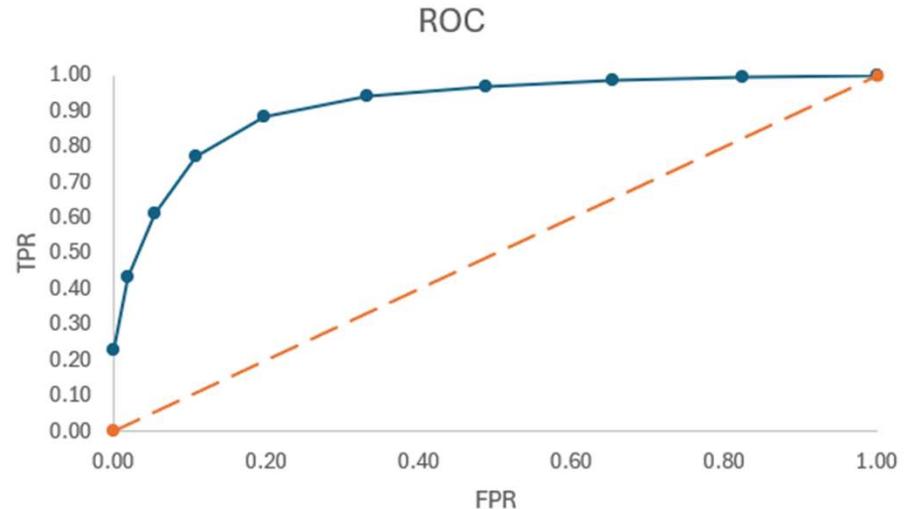
Ex.ID	Predicted	Actual	TPR	FPR
12	0.97	1	1/8 = 0.13	0/8 = 0.00
7	0.96	1	2/8 = 0.25	0/8 = 0.00
15	0.89	0	2/8 = 0.25	1/8 = 0.13
11	0.72	1	3/8 = 0.38	1/8 = 0.13
1	0.68	1	4/8 = 0.50	1/8 = 0.13
5	0.67	1	5/8 = 0.63	1/8 = 0.13
4	0.54	0	5/8 = 0.63	2/8 = 0.25
2	0.51	1	6/8 = 0.75	2/8 = 0.25
8	0.51	0	6/8 = 0.75	3/8 = 0.38
3	0.50	0	6/8 = 0.75	4/8 = 0.50
13	0.48	1	7/8 = 0.88	4/8 = 0.50
6	0.47	0	7/8 = 0.88	5/8 = 0.63
10	0.42	0	7/8 = 0.88	6/8 = 0.75
14	0.17	1	8/8 = 1.00	6/8 = 0.75
9	0.09	0	8/8 = 1.00	7/8 = 0.88
16	0.03	0	8/8 = 1.00	8/8 = 1.00



# ROC

## More Examples

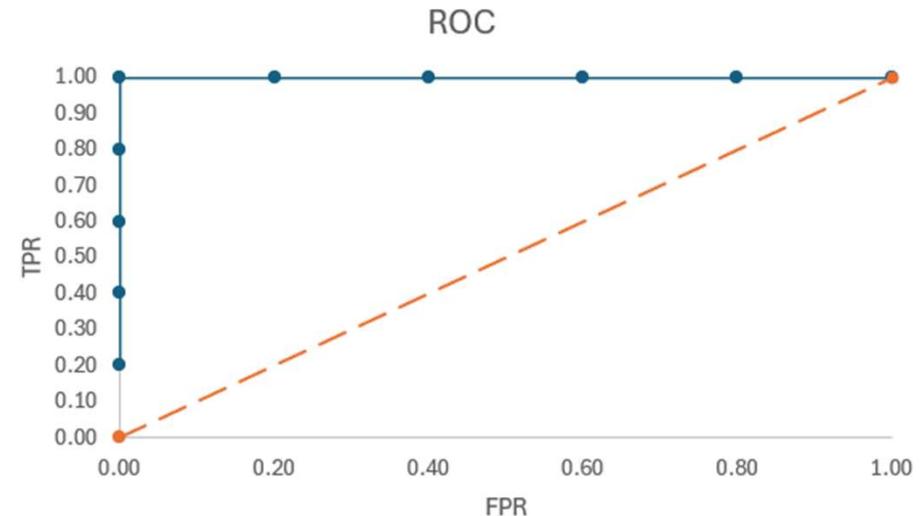
Prob(Y=1)	Distribution	count(True Y=1)	count(True Y=0)	TPR	FPR
[90%, 100%]	1,000	1,000	-	0.23	0.00
[80%, 90%)	1,000	900	100	0.43	0.02
[70%, 80%)	1,000	800	200	0.61	0.05
[60%, 70%)	1,000	700	300	0.77	0.11
[50%, 60%)	1,000	500	500	0.88	0.20
[40%, 50%)	1,000	250	750	0.94	0.33
[30%, 40%)	1,000	120	880	0.97	0.49
[20%, 30%)	1,000	80	920	0.99	0.65
[10%, 20%)	1,000	40	960	1.00	0.82
[0%, 10%)	1,000	20	980	1.00	1.00
Total	10,000	4,410	5,590		



# ROC

## More Examples

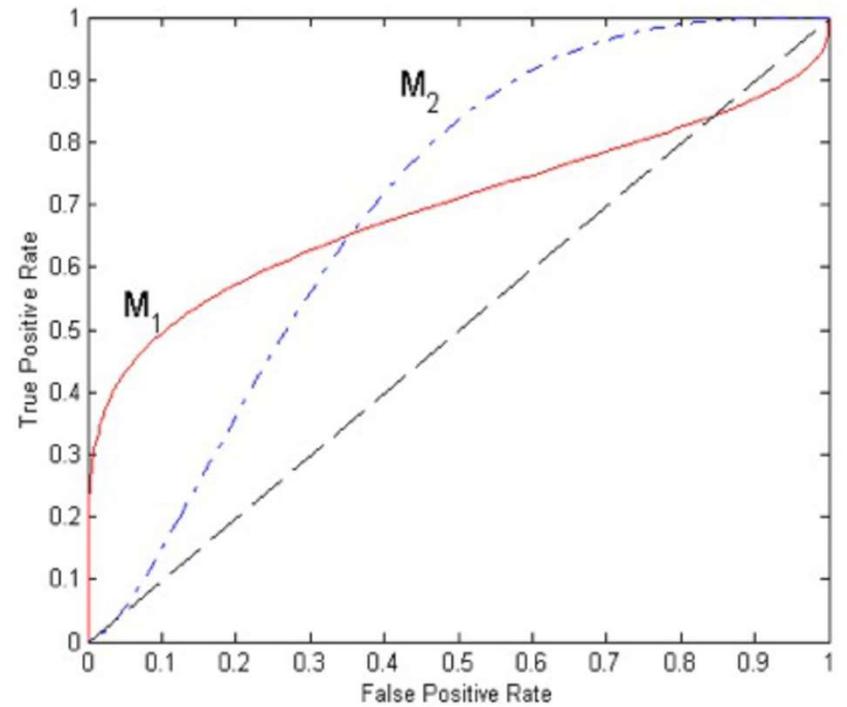
Prob(Y=1)	Distribution	count(True Y=1)	count(True Y=0)	TPR	FPR
[90%, 100%]	1,000	1,000	-	0.20	0.00
[80%, 90%)	1,000	1,000	-	0.40	0.00
[70%, 80%)	1,000	1,000	-	0.60	0.00
[60%, 70%)	1,000	1,000	-	0.80	0.00
[50%, 60%)	1,000	1,000	-	1.00	0.00
[40%, 50%)	1,000	-	1,000	1.00	0.20
[30%, 40%)	1,000	-	1,000	1.00	0.40
[20%, 30%)	1,000	-	1,000	1.00	0.60
[10%, 20%)	1,000	-	1,000	1.00	0.80
[0%, 10%)	1,000	-	1,000	1.00	1.00
Total	10,000	5,000	5,000		



# ROC

## ROC – Comparing Models

- Different algorithms can work better in different parts of ROC space.
- **M1** is better for small FPR
- **M2** is better for large FPR



# Area Under ROC Curve

## AUC – Area Under Curve

- Comprehensive measure: AUC evaluates model performance across all thresholds, not just one fixed decision threshold.
- Reflects the model's ability to distinguish between classes.
- Provides a single value for easy comparison of models.

