

---

# CS 4375 – Introduction to Machine Learning

---

Ensemble Methods – Boosting

Erick Parolin



[Based on the slides of Nicholas Ruozzi, Based on the slides of Vibhav Gogate and Rob Schapire]

# Previously

- **Variance reduction via bagging**
  - Generate “new” training data sets by sampling with replacement from the empirical distribution
  - Learn a classifier for each of the newly sampled sets
  - Combine the classifiers for prediction
- **How about reducing bias for binary classification problems?**

# Boosting

- How to translate ***rules of thumb*** into good learning algorithms?
- **Examples:**
  - In medical diagnosis, classifying patients as either at risk or not at risk for developing diabetes.
    - If the patient's body mass index (BMI) is above 30, classify them as at risk.
    - If the patient's age is over 45, classify them as at risk.
  - To classify emails as spam or not spam.
    - If the word "discount" appears in the email, classify it as spam.
    - If the sender's email address is unknown, classify it as spam.
  - To detect credit card fraud transactions
    - If the transaction amount is unusually high compared to the user's average spending, classify it as fraudulent.
    - If the transaction is made in a location far from the user's usual area, classify it as fraudulent.

# Boosting

- Freund & Schapire: Theory for “**weak learners**” in late 80’s
- **Weak Learner:** performance on any training set is slightly better than chance prediction
- Intended to answer a theoretical question, not as a practical way to improve learning
  - Tested in mid 90’s using *not-so-weak* learners
  - Works pretty well in practice!

# Boosting

- **Main idea:** use weak learner to create strong learner.
- **Ensemble method:** combine base classifiers returned by weak learner.
- Finding simple relatively accurate base classifiers often not hard.
- **But, how should base classifiers be combined?**

# Boosting

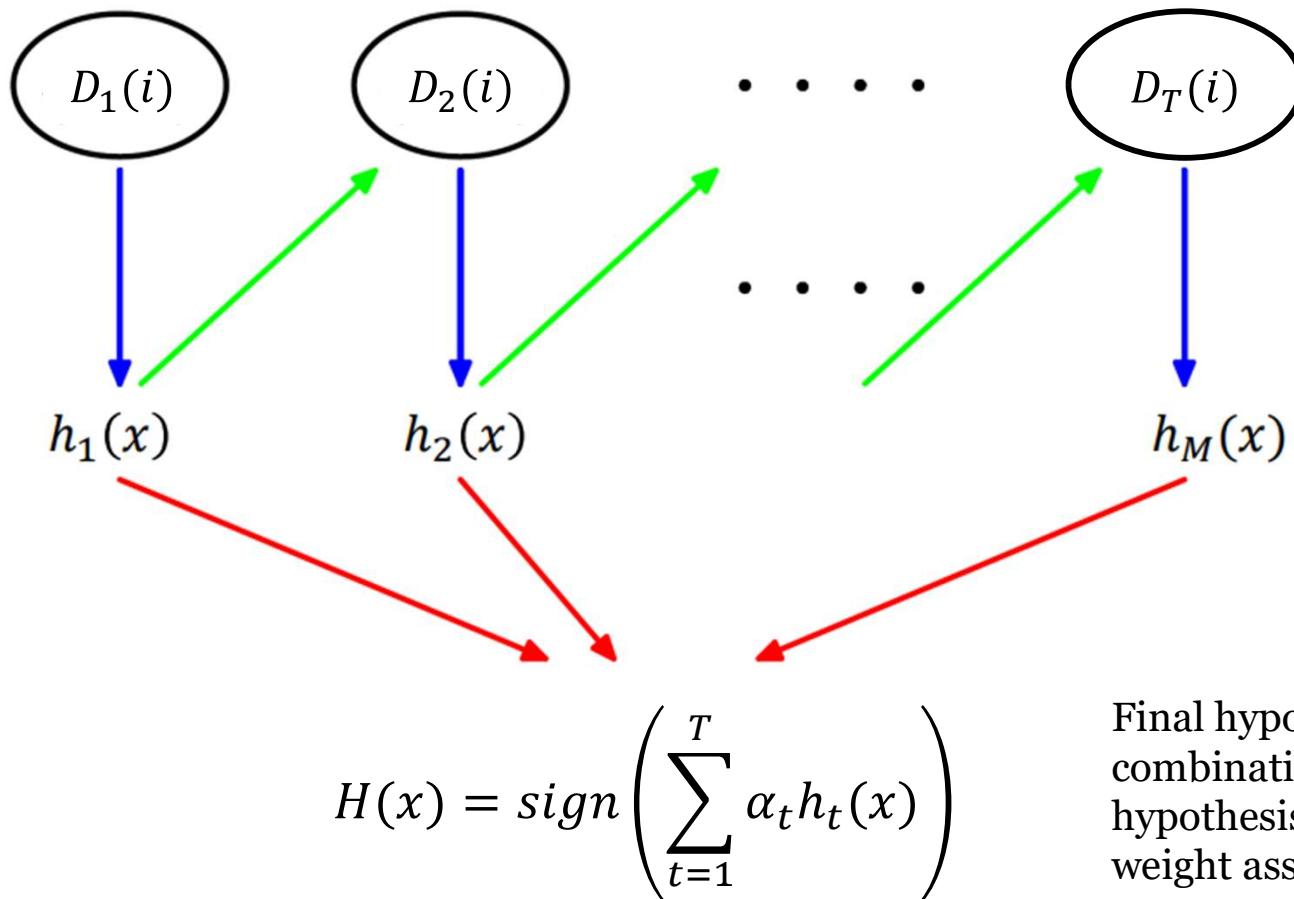
- Instead of learning a single (weak) classifier, learn many weak classifiers that are **good at different parts of the input space.**
- Output class: (Weighted) vote of each classifier
  - Classifiers that are most confident will vote with more conviction
  - Classifiers will be most confident about a particular part of the space
  - On average, it does better than single classifier!
    - **Force** classifiers to learn about different parts of the input space
    - Weight the votes of different classifiers

# Boosting – General Algorithm

- Weight all training samples equally
- Train model on training set
- Compute **error** of model on training set
- **Increase weights** on training cases model **gets wrong**
- Train new model on **re-weighted** training set
- Re-compute **errors** on weighted training set
- **Increase weights** again on cases model **gets wrong**  
**Repeat until tired (100+ iterations)**

**Final model:** weighted prediction of each model

# Boosting – Graphical Illustration



$D_t$  is a distribution that defines the weights for datapoints on round t.

Weak learner  $h_t$  will learn based on  $D_t$ , which is updated based on the training errors observed from  $h_{t-1}$

Final hypothesis H is composed by the combination of votes from each weak hypothesis (where  $\alpha_t$  corresponds to the weight associated to hypothesis  $h_t$ ).

# Adaboost

**Given:**  $(x_1, y_1), \dots, (x_m, y_m)$  where  $x_i \in X, y_i \in \{-1, +1\}$

**Initialize:**  $D_1(i) = 1/m$  for  $i = 1, \dots, m$

**For**  $t = 1, \dots, T$ :

Train weak learner using distribution  $D_t$

Get weak hypothesis  $h_t: X \rightarrow \{-1, +1\}$

Aim: select  $h_t$  with low weighted error:

$$\varepsilon_t = \Pr_{i \sim D_t}[h_t(x_i) \neq y_i]$$

$$\text{Choose } \alpha_t = \frac{1}{2} \ln \left( \frac{1 - \varepsilon_t}{\varepsilon_t} \right)$$

Update, for  $i = 1, \dots, m$ :

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

where  $Z_t$  is a normalization factor so that  $D_{t+1}$  sums to 1 ( $D_{t+1}$  will be a distribution)

**Output** the final hypothesis:  $H(x) = \text{sign}(\sum_{t=1}^T \alpha_t h_t(x))$

# Adaboost

**Given:**  $(x_1, y_1), \dots, (x_m, y_m)$  where  $x_i \in X, y_i \in \{-1, +1\}$

**Initialize:**  $D_1(i) = 1/m$  for  $i = 1, \dots, m$

$D_t(i)$  corresponds to the weight associated to  $i^{\text{th}}$  datapoint on  $t^{\text{th}}$  round/weak learner

**For**  $t = 1, \dots, T$ :

Train weak learner using distribution  $D_t$

Get weak hypothesis  $h_t: X \rightarrow \{-1, +1\}$

Aim: select  $h_t$  with low weighted error:

$$\varepsilon_t = \Pr_{i \sim D_t}[h_t(x_i) \neq y_i]$$

$$\text{Choose } \alpha_t = \frac{1}{2} \ln \left( \frac{1 - \varepsilon_t}{\varepsilon_t} \right)$$

Update, for  $i = 1, \dots, m$ :

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

**Weighted number of incorrect classifications of the  $t^{\text{th}}$  classifier.**

$$\varepsilon_t = \frac{\sum_i^m D_t(i) I(h_t(x_i) \neq y_i)}{\sum_i^m D_t(i)}$$

where  $I(\cdot)$  is an indicator function (1 if incorrect, 0 if correct).

where  $Z_t$  is a normalization factor so that  $D_{t+1}$  sums to 1 ( $D_{t+1}$  will be a distribution)

**Output** the final hypothesis:  $H(x) = \text{sign}(\sum_{t=1}^T \alpha_t h_t(x))$

# Adaboost

**Given:**  $(x_1, y_1), \dots, (x_m, y_m)$  where  $x_i \in X, y_i \in \{-1, +1\}$

**Initialize:**  $D_1(i) = 1/m$  for  $i = 1, \dots, m$

**For**  $t = 1, \dots, T$ :

Train weak learner using distribution  $D_t$

Get weak hypothesis  $h_t: X \rightarrow \{-1, +1\}$

Aim: select  $h_t$  with low weighted error:

$$\varepsilon_t = \Pr_{i \sim D_t}[h_t(x_i) \neq y_i]$$

$$\text{Choose } \alpha_t = \frac{1}{2} \ln \left( \frac{1 - \varepsilon_t}{\varepsilon_t} \right)$$

Update, for  $i = 1, \dots, m$ :

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

**$\alpha_t$  reflects the confidence in each weak learner  $t$ ,** better-performing learners get more weight, while worse-performing learners contribute less.

$$\varepsilon_t \rightarrow 0 \Rightarrow \alpha_t \rightarrow \infty$$

$$\varepsilon_t \rightarrow 0.5 \Rightarrow \alpha_t \rightarrow 0$$

$$\varepsilon_t \rightarrow 1 \Rightarrow \alpha_t \rightarrow -\infty$$

**Note:** If  $\varepsilon_t > 50\%$ , the weights are reverted back to  $1/m$ , and the resampling procedure is repeated.

where  $Z_t$  is a normalization factor so that  $D_{t+1}$  sums to 1 ( $D_{t+1}$  will be a distribution)

**Output** the final hypothesis:  $H(x) = \text{sign}(\sum_{t=1}^T \alpha_t h_t(x))$

# Adaboost

**Given:**  $(x_1, y_1), \dots, (x_m, y_m)$  where  $x_i \in X, y_i \in \{-1, +1\}$

**Initialize:**  $D_1(i) = 1/m$  for  $i = 1, \dots, m$

**For**  $t = 1, \dots, T$ :

Train weak learner using distribution  $D_t$

Get weak hypothesis  $h_t: X \rightarrow \{-1, +1\}$

Aim: select  $h_t$  with low weighted error:

$$\varepsilon_t = \Pr_{i \sim D_t}[h_t(x_i) \neq y_i]$$

$$\text{Choose } \alpha_t = \frac{1}{2} \ln \left( \frac{1 - \varepsilon_t}{\varepsilon_t} \right)$$

Update, for  $i = 1, \dots, m$ :

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

Updating the weights: if an example was misclassified by  $h_t$ , then  $D_{t+1}(i)$  increases for the next round, making  $i^{\text{th}}$  data point more important for the next weak learner.

$$\exp(-\alpha_t y_i h_t(x_i)) = e^{-\alpha_t} \text{ if } y_i = h_t(x_i)$$

$$\exp(-\alpha_t y_i h_t(x_i)) = e^{\alpha_t} \text{ if } y_i \neq h_t(x_i)$$

$$Z_t = \sum_{i=1}^m D_t(i) \exp(-\alpha_t y_i h_t(x_i))$$

where  $Z_t$  is a normalization factor so that  $D_{t+1}$  sums to 1 ( $D_{t+1}$  will be a distribution)

**Output** the final hypothesis:  $H(x) = \text{sign}(\sum_{t=1}^T \alpha_t h_t(x))$

# Adaboost

**Given:**  $(x_1, y_1), \dots, (x_m, y_m)$  where  $x_i \in X, y_i \in \{-1, +1\}$

**Initialize:**  $D_1(i) = 1/m$  for  $i = 1, \dots, m$

**For**  $t = 1, \dots, T$ :

Train weak learner using distribution  $D_t$

Get weak hypothesis  $h_t: X \rightarrow \{-1, +1\}$

Aim: select  $h_t$  with low weighted error:

$$\varepsilon_t = \Pr_{i \sim D_t}[h_t(x_i) \neq y_i]$$

$$\text{Choose } \alpha_t = \frac{1}{2} \ln \left( \frac{1 - \varepsilon_t}{\varepsilon_t} \right)$$

Update, for  $i = 1, \dots, m$ :

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

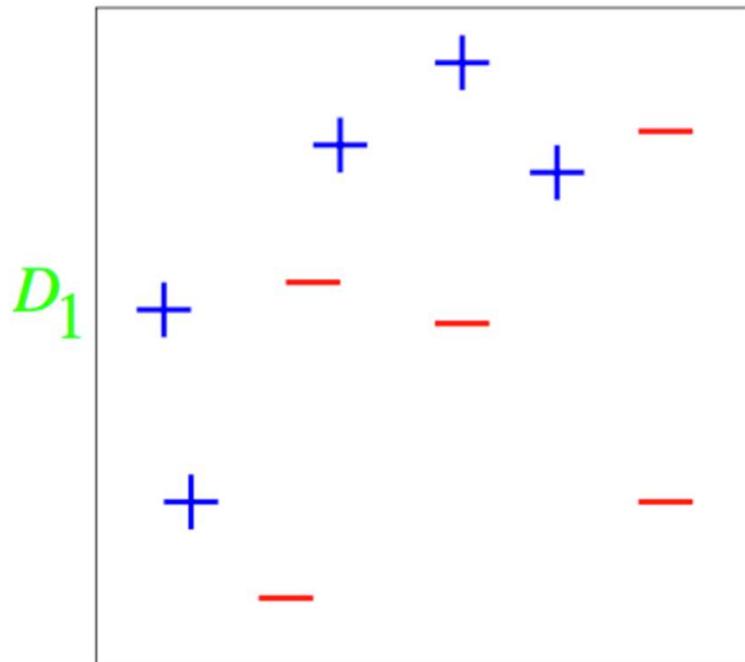
where  $Z_t$  is a normalization factor so that  $D_{t+1}$  sums to 1 ( $D_{t+1}$  will be a distribution)

**Output** the final hypothesis:  $H(x) = \text{sign}(\sum_{t=1}^T \alpha_t h_t(x))$

Voted combinations of each weak hypothesis, where the votes  $\alpha_t$  is used to emphasize component classifiers that are more reliable/confident than others.

# Adaboost – Numerical Example

Let's say we have the following training data: 10 data points, 5 of each class.



Then, we have  $m = 10$ , and following the algorithm, the initial weight distribution will be:

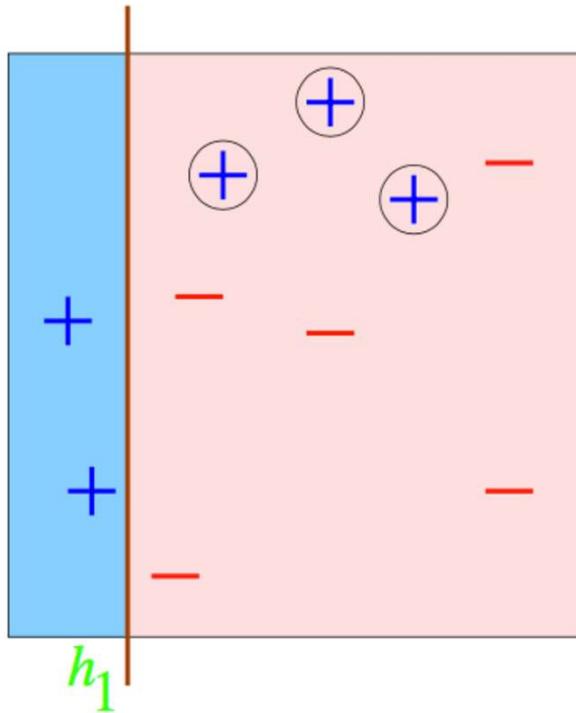
$$D_1 = 1/10$$

Which means:

$$D_1(1) = 0.1, D_1(2) = 0.1, \dots, D_1(10) = 0.1$$

# Adaboost – Numerical Example

Then, we train our first weak hypothesis  $h_1$



3 positive points are misclassified

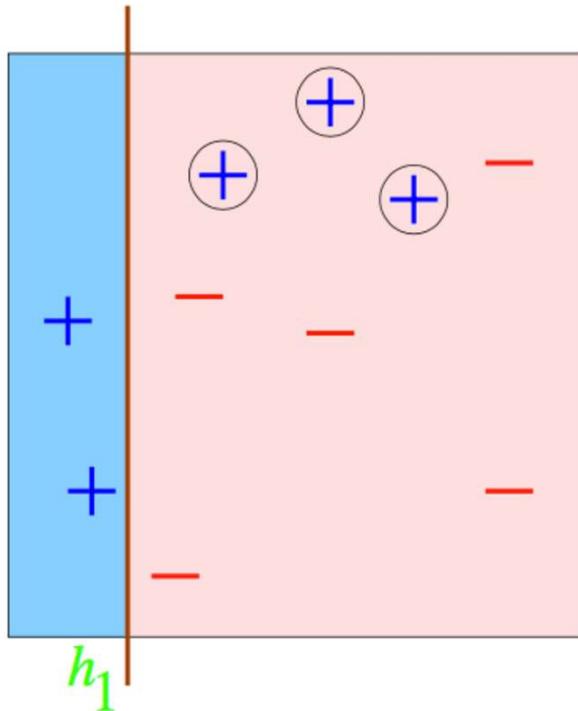
Then, we have:

$$\varepsilon_1 = \frac{\sum_{i=1}^{10} D_1(i) I(h_1(x_i) \neq y_i)}{\sum_{i=1}^m D_1(i)} = \frac{3 \times 0.1}{1} = 0.30$$

$$\alpha_1 = \frac{1}{2} \ln \left( \frac{1 - \varepsilon_1}{\varepsilon_1} \right) = \frac{1}{2} \ln \left( \frac{1 - 0.30}{0.30} \right) = 0.424$$

# Adaboost – Numerical Example

Computing new weight distribution  $D_2$



The weight points for the next round (round 2) is:

$$D_2(i) = \frac{D_1(i) \exp(-\alpha_1 y_i h_1(x_i))}{\sum_{i=1}^m D_1(i) \exp(-\alpha_1 y_i h_1(x_i))}$$

**Misclassified points:**  $0.1 \times \exp(0.424) = 0.153$

**Correctly classified points:**  $0.1 \times \exp(-0.424) = 0.065$

$$\sum_{i=1}^m D_1(i) \exp(-\alpha_1 y_i h_1(x_i)) = (3 \times 0.153) + (7 \times 0.065) = 0.914$$

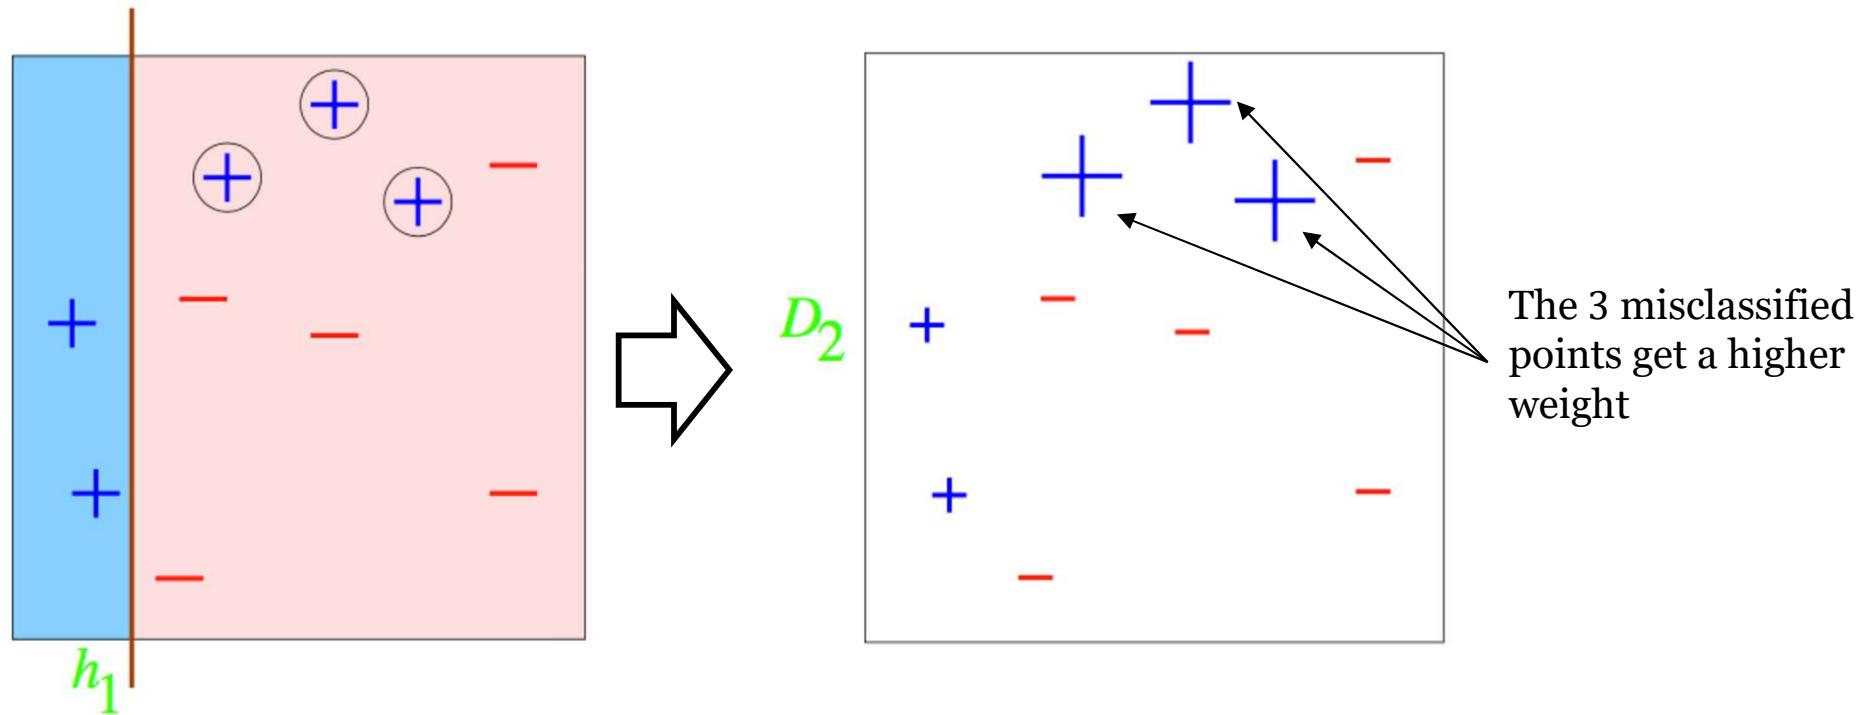
Normalizing, we have:

$$\text{Misclassified points: } D_2(i) = \frac{0.1 \times \exp(0.424)}{0.914} = 0.167$$

$$\text{Correctly classified points: } D_2(i) = \frac{0.1 \times \exp(-0.424)}{0.914} = 0.071$$

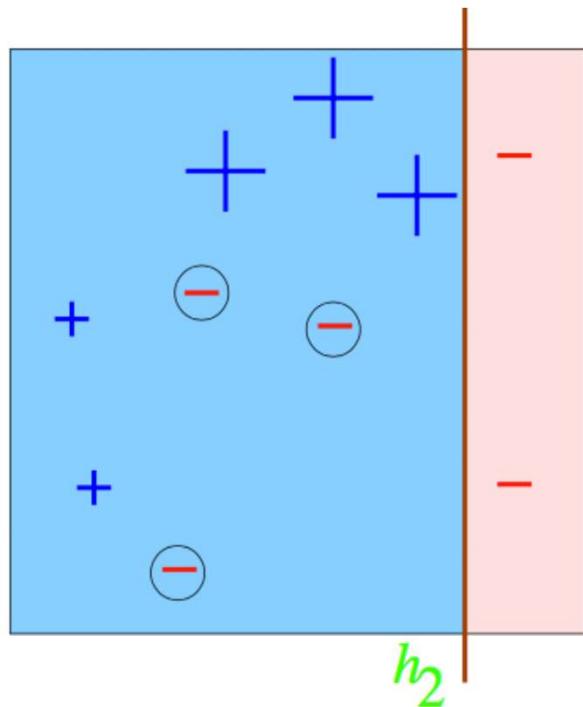
# Adaboost – Numerical Example

New weight distribution: After round 1, we get a new distribution  $D_2$



# Adaboost – Numerical Example

We train the weak hypothesis  $h_2$



3 positive points are misclassified

Then, we have:

$$\varepsilon_2 = \frac{\sum_{i=1}^{10} D_2(i) I(h_2(x_i) \neq y_i)}{\sum_{i=1}^m D_2(i)} = \frac{3 \times 0.071}{1} = 0.213$$

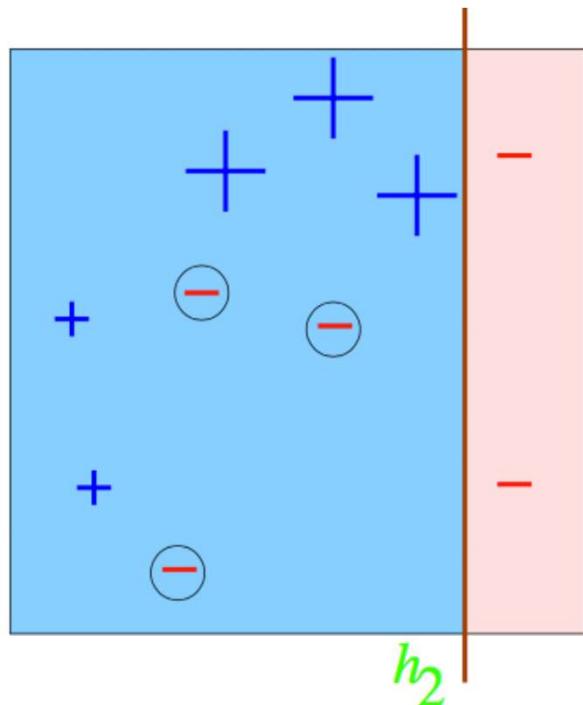
$$\alpha_2 = \frac{1}{2} \ln \left( \frac{1 - \varepsilon_2}{\varepsilon_2} \right) = \frac{1}{2} \ln \left( \frac{1 - 0.213}{0.213} \right) = 0.653$$

**Note:** we only used weight value  $D_2(i) = 0.071$  to compute  $\varepsilon_2$  because all the three misclassified points have the same weight (those points were not misclassified in round 1).

None of the points misclassified in round 1 were also misclassified in round 2.

# Adaboost – Numerical Example

Computing new weight distribution  $D_3$



The weight points for the next round (round 3) is:

$$D_3(i) = \frac{D_2(i) \exp(-\alpha_2 y_i h_2(x_i))}{\sum_{i=1}^m D_2(i) \exp(-\alpha_2 y_i h_2(x_i))}$$

**Misclassified points:**  $D_2(i) \exp(-\alpha_2 y_i h_2(x_i)) = 0.071 \times \exp(0.653) = 0.136$

**Correctly classified points:**

Larger +  $\Rightarrow D_2(i) \exp(-\alpha_2 y_i h_2(x_i)) = 0.167 \times \exp(-0.653) = 0.087$

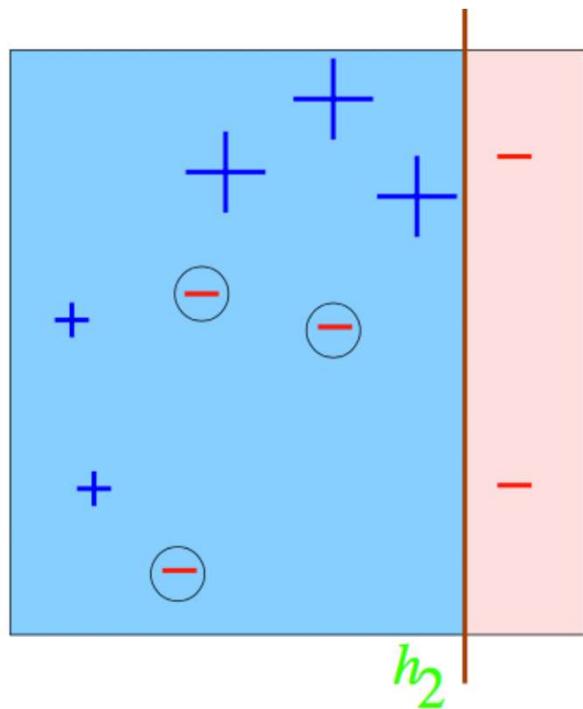
Smaller + and -  $\Rightarrow D_2(i) \exp(-\alpha_2 y_i h_2(x_i)) = 0.071 \times \exp(-0.653) = 0.037$

**Total Weight:**

$$Z_t = \sum_{i=1}^m D_2(i) \exp(-\alpha_2 y_i h_2(x_i)) = 3 \times 0.136 + 3 \times 0.087 + 4 \times 0.037 = 0.817$$

# Adaboost – Numerical Example

Computing new weight distribution  $D_3$



Normalizing, we have:

**Misclassified points:**  $\frac{0.136}{0.817} = 0.167$

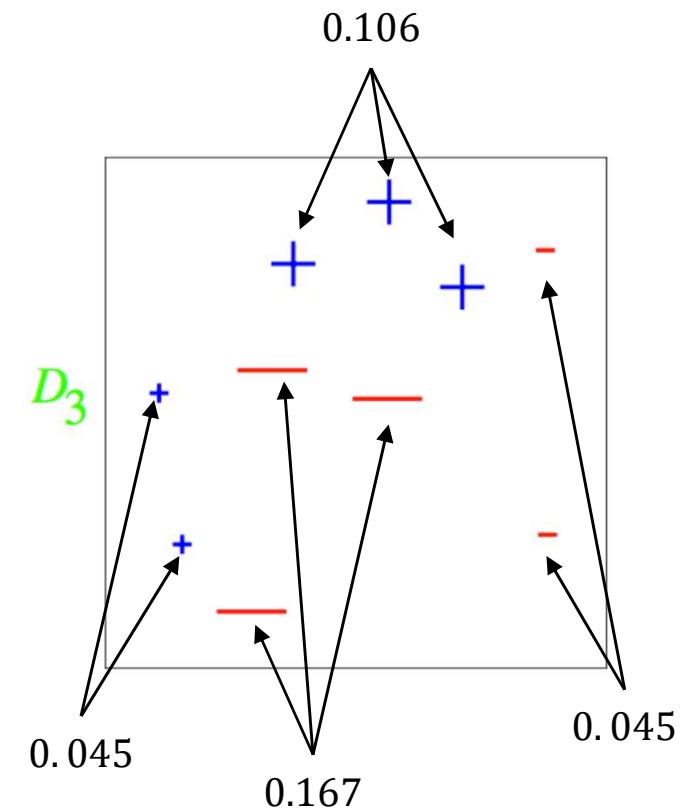
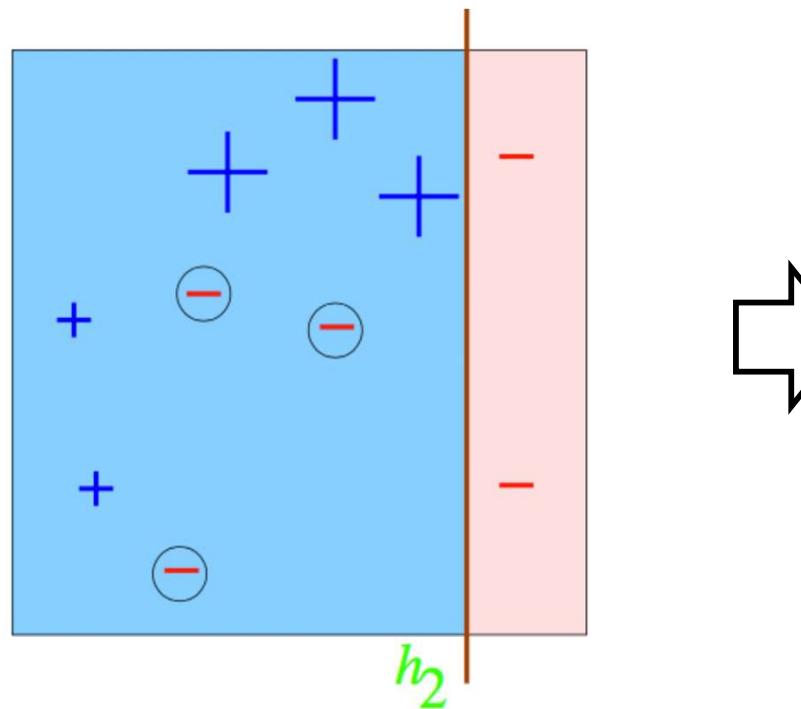
**Correctly classified points:**

Larger +  $\Rightarrow D_3(i) = \frac{0.087}{0.817} = 0.106$

Smaller +  $\Rightarrow D_3(i) = \frac{0.037}{0.817} = 0.045$

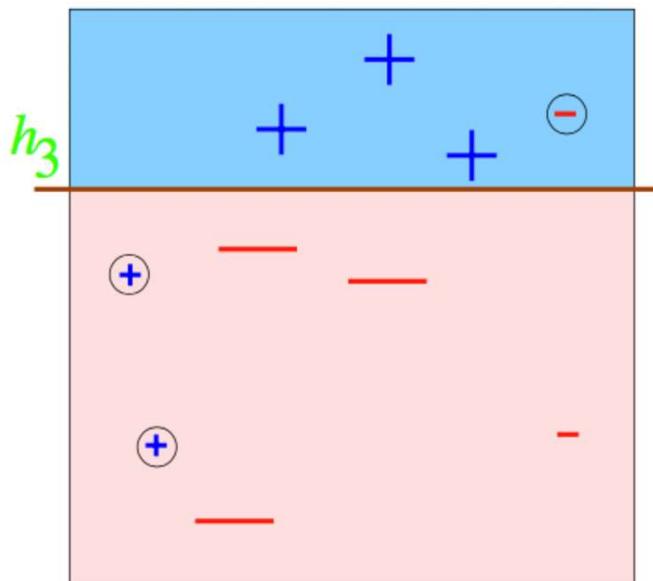
# Adaboost – Numerical Example

New weight distribution  $D_3$



# Adaboost – Numerical Example

Again, we train the weak hypothesis  $h_3$



3 positive points are misclassified

Then, we have:

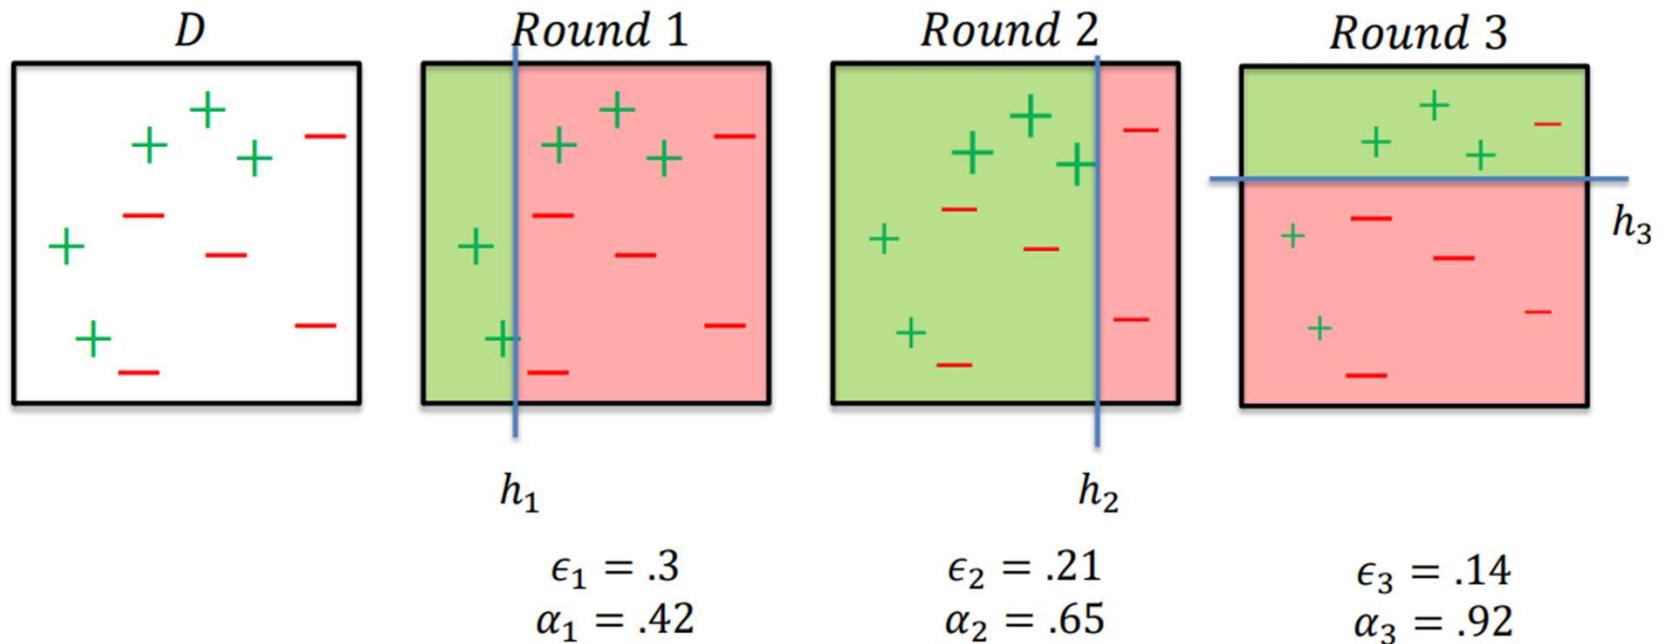
$$\varepsilon_3 = \frac{\sum_{i=1}^{10} D_3(i) I(h_2(x_i) \neq y_i)}{\sum_{i=1}^m D_3(i)} = \frac{3 \times 0.045}{1} = 0.135$$

$$\alpha_3 = \frac{1}{2} \ln \left( \frac{1 - \varepsilon_3}{\varepsilon_3} \right) = \frac{1}{2} \ln \left( \frac{1 - 0.135}{0.135} \right) = 0.929$$

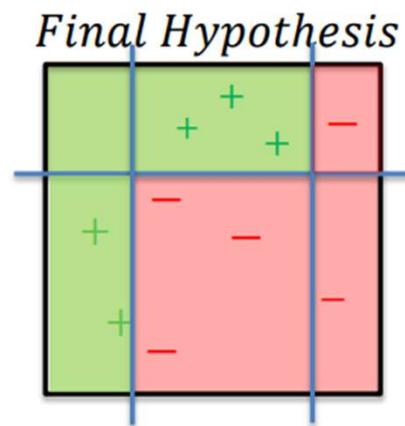
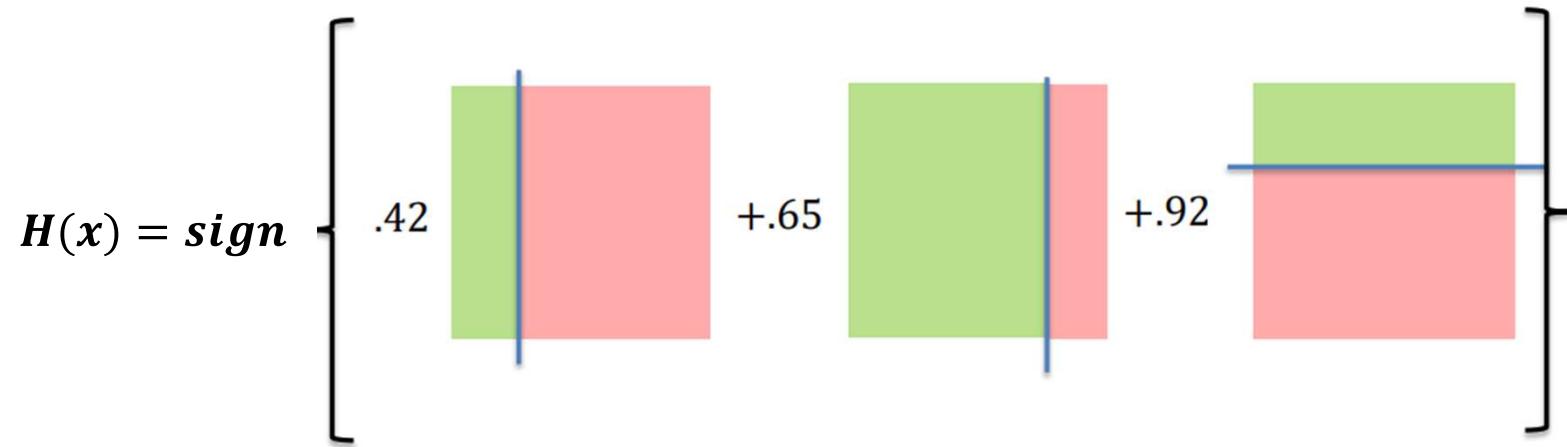
**Note:** again, we only used weight value  $D_3(i) = 0.045$  to compute  $\varepsilon_3$  because all the three misclassified points have the same weight (those points were not misclassified neither in round 1 nor in round 2).

# Adaboost – Numerical Example

In summary, that's what we computed along the three rounds



# Adaboost – Final Hypothesis



# Reweighting vs Resampling

- Example weights might be harder to deal with
  - Some learning methods can't use weights on examples
- We can resample instead:
  - Draw a bootstrap sample from the data with the probability of drawing each example proportional to its weight
- Reweighting usually works better but resampling is easier to implement

# **Summary: Boosting vs. Bagging**

- Bagging doesn't work so well with stable models. Boosting might still help.
- Boosting might hurt performance on noisy datasets. Bagging doesn't have this problem.
- On average, boosting helps more than bagging, but it is also more common for boosting to hurt performance.
- Bagging is easier to parallelize.

# **Readings**

- **Pattern Recognition and Machine Learning by Christopher M. Bishop – Chapter 14**
- **Explaining AdaBoost by Robert E. Schapire**